

AI AGENTS

The Future of Intelligent Automation

PRESENTER SCRIPT — 3 SPEAKERS

Full 60-Minute Session | Slides 1–24

SPEAKER	NAME PLACEHOLDER	SLIDES	APPROX. TIME
Person 1	Speaker 1	Slides 1–11	~20 minutes
Person 2	Speaker 2	Slides 12–18	~20 minutes
Person 3	Speaker 3	Slides 19–24	~20 minutes

COLOR LEGEND: Person 1 (Blue) Person 2 (Green) Person 3 (Purple) Audience Q (Orange)

Note to all speakers: Speak slowly and clearly. Pause after each key point. Make eye contact with the audience. You do NOT need to read every word — use this as a guide and speak naturally.

PERSON 1 — SLIDES 1 to 11 | Introduction, What AI Agents Are & How They Work

SLIDE 1 — Title Slide: AI AGENTS — The Future of Intelligent Automation

Person 1:

Good [morning/afternoon] everyone! Welcome, and thank you so much for being here today. We are really excited about what we're going to share with you in the next hour.

My name is [Your Name], and I'll be kicking things off. My two colleagues, [Name 2] and [Name 3], will take over later in the session.

Today we are going to talk about something that is changing the way work gets done — AI Agents. By the end of this session, you'll understand what they are, how they actually work behind the scenes, where they are already being used, and what the future looks like.

Don't worry — we're going to keep it simple and fun. No complicated technical jargon. Think of this as a friendly conversation about one of the most exciting things happening in tech right now.

Alright — let's dive in!

SLIDE 2 — Overview: What We'll Cover

Person 1:

Before we go deeper, here's a quick map of where we are headed today.

We'll start by understanding what an AI Agent actually is. Then we'll look at how it works inside — what happens step by step when an agent gets a task. After that, we'll explore different types of agents, where they're being used in the real world, what benefits they bring, and what challenges still exist. And finally, we'll look at where this technology is heading.

So it's a full journey — from basics to the future. Each section builds on the last, so by the end it will all come together nicely.

SLIDE 3 — Definition: What is an AI Agent?

Person 1:

Okay, so let's start with the big question — what exactly is an AI Agent?

Simply put, an AI Agent is a software program that can work on its own. You give it a goal, and it figures out how to reach that goal by itself. It doesn't need you to tell it every single step. It looks at the situation, thinks about what to do, takes action, and even learns from what happened.

Let me give you a real-life comparison. Think of a GPS navigation app. You tell it your destination. You don't tell it which roads to take, whether to avoid the toll, or when to recalculate if there's traffic. It figures all of that out itself. That's very similar to how an AI Agent works — you set the destination, it handles the journey.

Or here's another one: imagine a self-checkout machine at a supermarket. It scans items, calculates totals, handles payment, and even alerts a staff member if something goes wrong — all without a cashier guiding each step. It's operating with a goal and acting to achieve it. That's the spirit of an AI Agent.

Key things an AI Agent does: it takes in information from its environment, it reasons and makes a plan, it takes actions using tools available to it, and it learns from the results.

AUDIENCE QUESTION: Before we go on — quick question for the room: How many of you have used a voice assistant like Siri, Alexa, or Google Assistant? [Pause for hands.] Great! Those are actually early examples of AI Agents. They listen, they understand your request, and they take action. Today's AI Agents are much, much more powerful versions of that.

SLIDE 4 — Comparison: AI Agent vs Regular AI

Person 1:

Now I want to clear up something that confuses a lot of people. What's the difference between something like ChatGPT and an actual AI Agent?

Regular AI — like the chatbot you've probably used — answers one question at a time. You ask, it answers, done. Then you ask the next thing. It doesn't remember your previous conversation once the session ends, it can't browse the internet on its own, and it waits for you to give it every instruction.

An AI Agent is completely different. Think of it this way: regular AI is like a calculator — brilliant at one thing, but it just sits there until you press buttons. An AI Agent is more like a pilot with a full cockpit. It monitors, it decides, it acts, it adjusts — all in real time without you pressing every button.

A regular AI is reactive. An AI Agent is proactive. It doesn't just answer — it goes and does things. It can send emails, run code, search the web, and remember everything it learned last time. That's the

key difference.

SLIDE 5 — How It Works: The Brain of an AI Agent (Perceive → Think → Act → Remember)

Person 1:

Alright, now let's look at what actually happens inside an AI Agent when it's working on a task.

There are four steps, and they happen in a loop.

Step one is PERCEIVE. The agent takes in information. This could be a message you sent, a document, data from a website, or even an image. It's basically reading the situation.

Step two is THINK. The agent uses an AI model — its brain — to figure out: what is the goal here? What steps do I need to take? What tools do I have? It makes a plan.

Step three is ACT. It actually does something. It might send an email, search the web, run a piece of code, or call an external service. It executes the plan.

Step four is REMEMBER. It stores what happened — what worked, what didn't — so it can do better next time.

And then the loop starts again. Perceive, Think, Act, Remember — over and over until the job is done.

A good analogy here: think of a fire alarm system in a building. It continuously senses heat and smoke (Perceive), checks if the reading crosses a danger threshold (Think), triggers the sprinklers and calls emergency services (Act), and logs the event for maintenance review (Remember). Simple, but it's that same loop running non-stop.

AUDIENCE QUESTION: Quick check-in: Does this four-step loop make sense? Can anyone think of something in daily life that follows a similar pattern — sense, decide, act, learn? [Pause. Accept answers — traffic signals, thermostats, even humans learning on the job are all great answers.]

SLIDE 6 — Types: Types of AI Agents

Person 1:

Now that we know how agents work in general, let me quickly walk you through the different types. Think of it like different levels of ability — from basic to advanced.

The simplest is a Simple Reflex Agent. It just follows fixed rules. If this happens, do that. No memory, no planning. A great example is a spam filter in your email inbox. If the email says 'click here to win a prize,' it goes to spam. Simple rule, instant action.

Next is a Model-Based Agent. This one has memory. It tracks what happened before to make smarter decisions. Like a hospital patient monitoring system — it doesn't just look at your current blood pressure reading, it compares it against your history to decide if something is wrong.

Then we have Goal-Based Agents. These plan multiple steps to reach an objective. Like a warehouse robot that needs to pick up Box A, navigate around shelves, avoid other robots, and deliver to the packing station — it's planning a whole sequence.

Utility-Based Agents go further — they don't just reach the goal, they find the BEST way. Think of a food delivery app's routing system. It doesn't just find any route to deliver your food — it picks the fastest one, accounting for traffic and distance, to maximise customer satisfaction.

Learning Agents improve over time. The more they work, the smarter they get. And finally, Multi-Agent Systems are teams of agents — each one handling a specialised job, together solving something bigger. We'll come back to that in a bit.

SLIDE 7 — Components: Key Components of an AI Agent

Person 1:

Let's take a quick look at the parts that make up an AI Agent — like looking under the bonnet of a car.

At the centre is the LLM Brain — this is the Large Language Model, like GPT or Claude. It's the thinking engine. Without this, the agent can't understand language or make complex decisions.

Then there's Perception — the ability to receive and understand input. Memory — the ability to store information across tasks. Tools — these are things the agent can use to interact with the world, like sending emails or running code. And Planning — the ability to break a big goal into smaller manageable steps.

All five parts work together. Remove any one of them and the agent becomes much less capable. Together, they create something that can genuinely take on complex work.

SLIDE 8 — Applications: Where AI Agents Are Used Today

Person 1:

So where are AI Agents actually being used right now? Let me take you through some real examples — and these are not future plans, these are happening today.

Customer Support: Banks and e-commerce companies use AI agents that run 24/7. They understand what the customer needs, check order status, process refunds, and hand off to a human only when truly necessary. Imagine calling your bank at 2am and getting instant help — that's this.

Coding Assistants: Tools like GitHub Copilot don't just suggest lines of code — they understand the whole project and write entire functions, spot bugs, and test the code automatically.

Finance and Trading: In stock markets, AI agents analyse millions of data points per second and execute trades in milliseconds — far faster than any human trader could.

Research: Instead of a researcher spending three days reading papers, an AI research agent can browse hundreds of documents overnight and produce a structured summary by morning.

Travel Planning: Agents that simultaneously talk to airlines, hotels, and car rental services to build you a complete trip — in seconds.

Cybersecurity: These agents watch your network 24/7. The moment they detect something suspicious, they act — blocking the threat before a human even sees the alert.

The common theme? These agents are doing real work, in the real world, right now — not in a lab.

AUDIENCE QUESTION: I'm curious — looking at these applications, which one surprises you the most? And can anyone think of a task in your own job or life that an AI Agent could potentially take over? [Pause for 2-3 responses — this is a great engagement moment.]

SLIDE 9 — Data: AI Agent Adoption by Industry (Gartner Chart)

Person 1:

This chart from Gartner's 2024-2025 report shows something really striking. The industries adopting AI Agents fastest are not just tech companies. Healthcare, retail, financial services, manufacturing — they are all in.

The takeaway here is simple: this is not a niche technology for Silicon Valley startups. This is a mainstream shift happening across every major industry. If you work in any of these sectors, AI Agents are coming to your world — if they haven't arrived already.

SLIDE 10 — Tools: AI Agents Use Tools Like Humans

Person 1:

Here's something that I find really intuitive once you hear it: AI Agents use tools, just like we do.

Think about how you work. You use Google to search. You use your phone to send messages. You use Excel for data. You use your email for communication. You don't do everything in your head — you use tools.

AI Agents are the same. They have access to a toolkit: web search, code execution, database access, APIs that connect to other services, email and chat systems, file storage. The agent decides which tool to use and when — based on what the task needs.

A nice way to picture it: imagine a professional chef. They have knives, pans, ovens, blenders — each for a different job. The chef's knowledge tells them which tool to pick for which dish. The AI Agent's brain does the same thing — picks the right tool for the right task.

SLIDE 11 — Examples: Popular AI Agents You Know

Person 1:

Before I hand over to my colleague, let's look at some real AI Agents that are already out there in the world.

AutoGPT was one of the first agents the public could try. You give it a goal — say, 'research the top 5 electric cars and write a comparison report' — and it goes off and does all of that by itself.

Devin is described as the world's first AI software engineer. A startup showed it a brief description of an app, and Devin built the entire thing — frontend, backend, deployment — without a human writing a single line of code.

Microsoft Copilot inside Office 365 is something many of you might actually have access to. It drafts emails, creates slides from notes, summarises long documents, and analyses data in Excel — all as an agent inside your existing tools.

GPT Actions connects ChatGPT to apps like Gmail, Slack, and databases — so you can say 'check my emails, find the latest invoice from Vendor X, and update the spreadsheet' — and it does all three steps automatically.

These are not prototypes. These are live products that real teams are using today.

Alright — that covers the foundations. I'm now going to hand over to [Name 2] who will take us deeper into how these agents work together and why they're so powerful.

PERSON 2 — SLIDES 12 to 18 | Multi-Agent Systems, Benefits, Challenges & How to Build

SLIDE 12 — Advanced: Multi-Agent Systems: Teamwork!

Person 2:

Thank you, [Name 1]! Hello everyone, I'm [Your Name] and I'll be covering the next part of our session.

We've learned what a single AI Agent is. Now let's talk about what happens when you put multiple agents together — because that's where things get really interesting.

Think about a hospital emergency department. When a patient comes in, it's not one person handling everything. There's a triage nurse, a doctor, a radiologist, a pharmacist, and an administrator — all working in parallel, each doing their specialised job, and communicating with each other. Together, they solve something that no single person could do alone as quickly.

Multi-Agent Systems work exactly like that. You have a Manager Agent that receives the overall task and breaks it down. Then it delegates to specialised agents: a Research Agent that gathers information, a Writing Agent that creates content, a Coder Agent that builds something, and an Email Agent that handles communication. They all run in parallel and report back.

The slide shows an example: someone asks for a 'marketing campaign.' The Manager Agent splits that into research, content writing, and email delivery — three things happening simultaneously. What might take a team of humans a full day happens in minutes.

This is why multi-agent systems are so exciting. The speed isn't just about one agent being fast. It's about many agents running at the same time, each doing what it's best at.

AUDIENCE QUESTION: Here's a fun thought experiment: If you could build a team of AI agents for your own team at work — what roles would you give them? What would your 'Manager Agent' assign to each one? [Allow 2-3 quick answers — this makes the concept personal and memorable.]

SLIDE 13 — Benefits: Why AI Agents Are Powerful

Person 2:

Now let's talk about why organisations are so excited about this technology. What does it actually deliver?

Speed: AI Agents can complete tasks 100 times faster than humans. A task that a team member might spend three hours on — gathering data, organising it, writing a report — an agent does in under two minutes. Not because it's cutting corners, but because it doesn't pause, doesn't context-switch, and doesn't need to check its phone.

Accuracy: 99.9% accuracy — no fatigue mistakes. When humans do repetitive work, errors creep in. AI Agents don't get tired. A quality check agent running at 11pm is just as accurate as it was at 9am.

Scale: This is perhaps the most mind-bending one. One AI Agent can handle a thousand tasks simultaneously. Think about that — one system, a thousand parallel workstreams.

Cost: Over time, automating repetitive tasks reduces operational costs by around 60%. The upfront investment pays for itself quickly.

Always On: No weekends, no holidays, no sick days. An AI Agent watching your cybersecurity at 3am on a Sunday is just as alert as it would be on a Tuesday afternoon.

Adaptability: And unlike traditional software that breaks when conditions change, AI Agents can learn and adjust. They improve with experience.

Together, these benefits mean that organisations can do things they simply could not afford to do before — at a speed and quality that was previously impossible.

SLIDE 14 — Challenges: Challenges & Limitations

Person 2:

Now, and this is really important — AI Agents are not perfect. Let's be honest about the challenges, because understanding the limitations helps you use the technology responsibly.

Hallucinations: This is the term for when an AI confidently gives you wrong information. An AI Agent might complete a research task but include a statistic that doesn't exist — and present it as fact. This is why human review matters, especially for important decisions.

Security Risks: An agent that has access to your email, files, and databases is also a target for attackers. There are techniques called 'prompt injection' where a bad actor tricks the agent into doing something harmful by hiding instructions in the data it reads. Think of it like a note inside a parcel that says 'deliver this box to a different address' — the agent just follows it.

Cost: Each AI decision involves calling a large language model, which has a cost. For small tasks it's fine, but at scale, poorly designed agents can run up significant bills. Good design matters.

Over-Automation Risk: If you let an agent run without checkpoints and it makes an error early on, that error can snowball through many automated steps before anyone notices. Human oversight at key points is essential.

Lack of Common Sense: AI Agents can solve very complex technical problems but can completely miss something obvious that any five-year-old would know. They don't have life experience or intuition.

Privacy: Agents that access personal data must be carefully governed. Who sees what the agent saw? How is data stored? These are compliance questions that teams must answer before deploying agents.

The message here isn't to be scared of the technology — it's to use it thoughtfully. With the right guardrails, these challenges are very manageable.

AUDIENCE QUESTION: Out of these six challenges, which one concerns you the most — personally or professionally? And do you have any ideas about how it could be managed? [This is a great discussion point — allow 3-4 responses. Security and hallucinations usually generate the most discussion.]

SLIDE 15 — Market Data: The AI Agent Market is Exploding

Person 2:

Let's put some numbers on this. The AI Agent market is projected to reach 102 billion dollars by 2028, growing at 45 percent per year.

To put that in context: the global smartphone market today is about 500 billion dollars — built over 15 years. AI Agents are on track to reach a quarter of that size in just four years. That's an extraordinary rate of adoption.

What's driving this? Cost efficiency is a big factor. But also, the technology has finally crossed a threshold where it actually works well enough for real business use. A few years ago, agents were unreliable and needed constant babysitting. Today, they're robust enough to handle real workloads.

If you're in a leadership role, or you're thinking about your career — this number tells you something important: organisations that figure out how to use AI Agents well will have a significant competitive advantage over those that don't. This is not a trend to wait and see on.

SLIDE 16 — Build: How to Build a Simple AI Agent

Person 2:

Now — for those of you who are curious about the 'how do you actually build one' — this slide walks through the six steps.

Step 1 is Define the Goal. Be very specific. 'Answer customer emails' is too vague. 'Respond to customer emails about order delivery within 2 hours, using our shipping database to give accurate updates' — that's a proper goal.

Step 2 is Choose the Brain. Which AI model powers the agent's thinking? OpenAI's GPT-4, Anthropic's Claude, Google's Gemini — these are all options. Each has different strengths.

Step 3 is Give It Tools. What can the agent actually do? Connect it to the systems it needs — a database, an email service, a web search API.

Step 4 is Add Memory. Should the agent remember past conversations? A customer support agent that doesn't remember what a customer said ten minutes ago is not very useful.

Step 5 is Set Safety Rules. What should it never do? What requires human approval? Every agent needs a rulebook.

Step 6 is Test and Improve. Run real scenarios. See where it fails. Fix it. Repeat.

You don't need to be a programmer to understand this process. It's the same logic as setting up a new member of staff — you define their role, give them the right tools, tell them the rules, and then check their work as they get started.

SLIDE 17 — Frameworks: Popular AI Agent Frameworks

Person 2:

For the more technical people in the room, this slide shows the main toolkits — called frameworks — that developers use to build AI Agents.

LangChain is the most popular starting point. It makes it easy to connect an LLM brain to memory and tools. Great for beginners.

LangGraph is an extension that handles more complex flows — where the agent doesn't just go in one direction but might loop back or branch based on what it finds.

AutoGen, from Microsoft, specialises in multi-agent setups — where multiple AI agents talk to each other and collaborate to solve problems.

CrewAI lets you define a team of agents with specific roles — like a Project Manager agent, a Researcher agent, and a Writer agent — each with their own focus.

Semantic Kernel is Microsoft's enterprise-grade option for large-scale production deployments.

If you're just curious and want to try something today, LangChain has excellent tutorials online. You can have a basic agent running in an afternoon.

And with that, I'll hand over to [Name 3] who will take us through the real-world case studies, the future, and close out our session.

PERSON 3 — SLIDES 18 to 24 | Future, LLMs, RAG, Case Studies & Closing

SLIDE 18 — Future: The Future of AI Agents (Timeline)

Person 3:

Thank you, [Name 2]! Hello everyone — I'm [Your Name] and I'll be taking you through the final stretch, including some amazing real-world examples and where all of this is heading.

Let's start with a timeline, because I think it helps to see the speed at which this technology is moving.

In 2024 — which is already behind us — agents became capable of browsing the web, running code, and reading files. That was the milestone year where agents started doing genuinely useful work.

In 2025 — right now — we're in the era of Agentic Workflows. Agents are handling long, multi-step tasks that run for hours or even days without human input. A business might kick off an agent on Monday morning and receive a completed research report by Wednesday.

By 2026 to 2027, physical AI Agents will be mainstream. Robots in warehouses, hospitals, and construction sites — controlled by AI agent brains. We already see early versions of this with companies like Boston Dynamics and Amazon's warehouse robots.

And by 2028 to 2030, the predictions are society-scale. Agents managing city infrastructure — traffic systems, energy grids, healthcare systems. These are not science fiction — they are extrapolations from what's already being built today.

Sam Altman, the CEO of OpenAI, said recently that AI Agents will do the work of a small company. Humans will move from doing tasks themselves to managing agents that do the tasks. That's a fundamental shift in what work looks like.

AUDIENCE QUESTION: Looking at this timeline — honestly, how does it make you feel? Excited? A little nervous? Both? [Pause — this is a genuine conversation starter. Most people feel a mix. Acknowledge all responses without judgment. Then continue.] Both reactions are completely valid. The key is to be informed and prepared — which is exactly why you're here today.

SLIDE 20 — LLM Brain: The Role of LLM in AI Agents

Person 3:

Let me take a moment to explain the brain behind all of this — the Large Language Model, or LLM.

An LLM is what makes the agent intelligent rather than just automated. Traditional automation is rule-based — if X happens, do Y. An LLM can understand nuance, context, and intent.

Here's a concrete comparison: a traditional automated system for customer emails might be programmed to respond to the word 'refund.' But what if a customer writes 'I've been waiting three weeks and I'm really not happy'? A rule-based system might not catch that. An LLM-powered agent

understands the frustration, recognises it as a refund or escalation situation, and responds appropriately.

The LLM is also what allows agents to use tools intelligently. It decides: 'For this step, I need to search the web. For this step, I need to query the database. For this step, I need to send an email.' It's the decision-maker that coordinates everything else.

Without an LLM, you have automation. With an LLM, you have intelligence. That's the difference.

SLIDE 21 — Knowledge Engine: RAG — Retrieval-Augmented Generation

Person 3:

Now here's a concept that sounds technical but is actually very logical once you understand it: RAG, which stands for Retrieval-Augmented Generation.

The problem with AI models is that they were trained on data up to a certain date. They don't automatically know about your company's internal policies, your products, or the latest news.

RAG solves this. Before the agent answers a question, it first searches a knowledge base — your company documents, a database, medical records, legal files — retrieves the relevant information, and then uses that as context when generating the answer.

Think of it like this: imagine you just joined a company and someone asks you a technical question about your product. You don't know the answer off the top of your head, but you know exactly where to look — the product manual. You grab it, read the relevant section, and then answer confidently and accurately. That's RAG. The agent reaches for the right document before speaking.

This is what prevents hallucinations in high-stakes situations. Instead of the agent guessing, it checks the source. Medical AI that accesses actual patient records. Legal AI that reads actual case files. Customer support that references actual product documentation.

RAG is what makes AI Agents reliable for professional, regulated industries.

AUDIENCE QUESTION: *Can anyone think of a situation in your own work where having an AI that could instantly search your company's internal documents — and give you a precise, accurate answer — would save you significant time? [Pause for 2-3 responses. This often leads to great examples around onboarding, compliance checks, or policy queries.]*

SLIDE 22 — Market Today: Netflix — AI Agents in Action

Person 3:

Let's look at a company most of you use: Netflix.

Netflix has 250 million subscribers across 190 countries. Every single one of them sees a different home screen when they log in. How is that possible? AI Agents.

Netflix's recommendation engine is an AI Agent that watches what you watch, what you pause on, what you search for, and what similar users with your taste patterns watch — and uses all of that to decide exactly which 40 titles to show you, in what order, with which thumbnail image.

And it's not just recommendations. AI agents at Netflix test video playback across 50 different device types automatically. They predict which users are likely to cancel their subscription next month — and trigger a personalised offer before they do. They optimise server infrastructure to reduce costs.

The results are staggering: 80 percent of everything you watch on Netflix was suggested by an AI Agent — not chosen by you from browsing. 40 percent lower churn rate than the industry average. Over a billion dollars saved per year in infrastructure costs alone.

This is not a future use case. This is live, at massive scale, delivering measurable business value every single day.

SLIDE 23 — Market Today: Tesla — AI Agents in Autonomous Vehicles

Person 3:

Now let's look at perhaps the most dramatic real-world deployment of AI Agents: Tesla's autonomous driving system.

When a Tesla is driving itself, it's not running one AI — it's running a coordinated team of AI Agents in real time.

The Vision Agent processes images from 8 cameras, analysing 2.5 billion frames every single day. It detects pedestrians, cyclists, lane markings, traffic signs — in real time, with 99.9% accuracy.

The Decision Agent takes all that visual information and makes split-second choices: should I brake? change lane? accelerate? It plans 5 to 10 seconds ahead at every moment.

The Safety Agent is monitoring the vehicle itself — temperature, battery, brake pressure. It'll alert you to a problem before it becomes an emergency.

And then there's the Fleet Learning Agent — my personal favourite. There are 1.5 million Teslas on the road. Every unusual driving situation encountered by any of them is sent back, anonymously, and used to teach all the other Teslas how to handle that situation better. It's like 1.5 million drivers sharing their experience every day. No human training programme could ever achieve that kind of scale.

This is AI Agents working together at a level of complexity and real-time precision that simply would not be possible any other way.

SLIDE 19 — Summary: Key Takeaways

Person 3:

Before we close, let me bring it all together with the key points from today.

First: AI Agents are software programs that can perceive, think, act, and learn — on their own — to achieve a goal. No step-by-step human instruction required.

Second: They use tools — web search, databases, email, APIs — and combine them intelligently to solve complex, multi-step problems.

Third: This isn't a niche technology. Every major industry is adopting it, right now.

Fourth: The benefits are real and measurable — speed, accuracy, scale, cost reduction, and 24/7 availability.

Fifth: The challenges are real too — hallucinations, security, cost, and privacy. Responsible deployment means having guardrails and human oversight.

Sixth: The market will reach 102 billion dollars by 2028. This is not a future thing. It is a now thing.

AUDIENCE QUESTION: OPEN Q&A; — We'd love to hear from you now. Any questions at all — about anything we covered today? Something you want us to go deeper on? Something you're unsure about? The floor is yours. [Allow 10-15 minutes for genuine Q&A.; All three presenters should feel free to answer.]

SLIDE 24 — Thank You & Closing

Person 3:

We are almost at time, so let me close with a few thoughts.

What we've talked about today is not just technology for technology's sake. AI Agents represent a genuine shift in how work gets done. The organisations and individuals who understand this early and start building with it will have a real advantage.

You don't need to become a developer to benefit from this. You just need to understand what's possible — and start asking the question: where in my work could an AI Agent help?

On the slide you'll see three recommendations: Keep Learning — follow AI developments and try tools. Start Building — even a simple experiment with LangChain or AutoGPT will teach you more than any presentation. Stay Updated — this field moves fast, and staying curious is the best skill you can have right now.

On behalf of myself, [Name 1], and [Name 2] — thank you so much for your time and your wonderful questions today. We hope this has been genuinely useful. Please come and talk to any of us afterwards if you want to continue the conversation.

Thank you!

End of Script | Total estimated duration: 55–65 minutes including Q&A pauses | Adjust pace to audience engagement — if discussion is rich, let it run.