AI AGENTS

The Future of Intelligent Automation

PRESENTER SCRIPT — 3 SPEAKERS

Full 60-Minute Session | Slides 1 – 24

SPEAKER	SLIDES	APPROX. TIME
Person 1	Slides 1 – 11	~20 minutes
Person 2	Slides 12 – 18	~20 minutes
Person 3	Slides 19 – 24	~20 minutes

Note to all speakers: Speak at a steady pace. Pause after each key point. Make eye contact with the audience. This script is a guide — feel free to speak naturally in your own words.

PERSON 1 — SLIDES 1 to 11 | Introduction · What AI Agents Are · How They Work

SLIDE 1 — Title Slide: AI AGENTS — The Future of Intelligent Automation

Person 1:
Good [morning / afternoon], everyone. Welcome, and thank you for joining us today.

My name is [Your Name]. I will be presenting the first part of our session. My colleagues [Name 2] and [Name 3] will take over from me later.

Over the next hour, we are going to talk about AI Agents — one of the most significant developments in technology right now. By the end of this session, you will understand what AI Agents are, how they work, where they are already being used, and where they are headed.

We will keep the language clear and straightforward. You do not need a technical background to follow along. Let us get started.

SLIDE 2 — Overview: What We Will Cover

Person 1:
Here is a quick overview of what we will cover today.

We start with the basics — what an AI Agent is and how it differs from regular AI. Then we look at the different types of agents and where they are used in the real world. After that, my colleagues will cover the benefits, the challenges, how agents are built, and what the future looks like.

Each section connects to the next, so by the end, you will have a complete picture.

SLIDE 3 — Definition: What is an AI Agent?

Person 1:
Let us start with the core question: what exactly is an AI Agent?

An AI Agent is a software program that can work on its own to complete a goal. You give it an objective, and it decides how to reach that objective by itself. It does not need you to spell out every step. It reads the situation, makes a plan, takes action, and learns from the outcome.

A useful comparison is a GPS navigation app. You enter your destination. You do not tell it which roads to take, whether to avoid the motorway, or what to do when there is traffic. The app figures all of that out. It perceives the road conditions, decides the best route, guides you turn by turn, and adjusts if you miss a junction. That is the same logic an AI Agent follows — you set the destination, and it manages the journey.

Another example: a self-checkout machine at a supermarket. It scans items, calculates totals, processes payment, and alerts a staff member if there is a problem — all without a cashier directing each step. It has a goal and it acts to achieve that goal.

To summarise: an AI Agent perceives its environment, reasons about the best course of action, takes action using available tools, and learns from the results.

AUDIENCE QUESTION Before we go on — a quick show of hands: how many of you have used a voice assistant such as Siri, Google Assistant, or Alexa? [Pause for hands.] Good. Those are early, basic examples of AI Agents. They listen, they interpret your request, and they take action. Today's AI Agents are far more capable versions of that idea.

SLIDE 4 — Comparison: AI Agent vs Regular AI

Person 1:
This is a question that comes up often: how is an AI Agent different from something like ChatGPT?

Regular AI responds to one prompt at a time. You ask a question, it answers, and then it waits for the next question. It has no memory of your previous conversation once the session ends. It cannot go off and do things on its own. It is entirely dependent on you for every instruction.

An AI Agent operates differently. Think of a regular AI like a reference book — incredibly useful, but it sits on the shelf until you pick it up and ask it something. An AI Agent is more like an air traffic controller — it monitors live information continuously, makes decisions in real time, coordinates multiple actions at once, and manages the full process from start to finish.

Regular AI is reactive. It responds when prompted. An AI Agent is proactive. It takes initiative, executes tasks, and works through problems without waiting for you to guide each step.

SLIDE 5 — How It Works: Perceive → Think → Act → Remember

Person 1:
Now let us look at what actually happens inside an AI Agent when it is working on a task.

The process follows four steps, and these steps repeat in a loop until the job is done.

Step 1 — Perceive. The agent takes in information. This could be a message, a document, data from a website, or an image. It reads the situation.

Step 2 — Think. The agent uses its AI brain to make a plan. It asks: what is the goal? What steps are needed? Which tools should I use?

Step 3 — Act. The agent executes the plan. It might search the web, send an email, run a piece of code, or retrieve data from a database.

Step 4 — Remember. The agent stores what it learned — what worked, what did not — so it can make better decisions next time.

A helpful analogy is a building's fire safety system. It continuously monitors heat and smoke levels (Perceive). When a reading crosses the danger threshold, it identifies the risk (Think). It activates the sprinklers and contacts emergency services (Act). It logs the event for maintenance records (Remember). That same loop runs inside every AI Agent.

AUDIENCE QUESTION Does this four-step loop make sense? Can anyone think of something in everyday life that works in a similar way — sense something, decide what to do, act, and then learn from it? [Pause. Good answers include traffic signals, thermostats, or even how a person learns a new skill on the job.]

SLIDE 6 — Types: Types of AI Agents

Person 1:
AI Agents come in different levels of capability. Let me walk through the main types.

Simple Reflex Agent. This is the most basic type. It follows fixed rules: if this happens, do that. No memory, no planning. A good example is an email spam filter. If an email contains phrases like 'limited time offer' or 'click to claim,' it is moved to the spam folder immediately, based on a rule.

Model-Based Agent. This type keeps a record of past events and uses that history to make better decisions. Consider a patient monitoring system in a hospital. It does not just look at your current blood pressure — it compares it against your previous readings to determine whether the change is significant.

Goal-Based Agent. This agent plans a sequence of steps to reach a specific objective. A warehouse robot is a good example. It needs to locate an item, navigate around obstacles, avoid other robots, and deliver the package to the correct station. It plans the whole sequence.

Utility-Based Agent. This goes a step further — it does not just reach the goal, it finds the best way to reach it. A food delivery app's routing system does not just find any route to your address — it finds the fastest one, accounting for traffic, distance, and delivery sequence.

Learning Agent. This type improves with every task it completes. The more experience it has, the smarter it becomes.

Multi-Agent System. This is a team of agents, each specialising in one area, working together to solve a complex problem. We will look at this in more detail shortly.

SLIDE 7 — Components: Key Components of an AI Agent

Person 1:
Let us briefly look at what an AI Agent is made of.

At the centre is the **LLM Brain** — a Large Language Model such as GPT or Claude. This is the thinking engine. It understands language, reasons through problems, and makes decisions.

Around it are four supporting components: **Perception** — the ability to receive and understand input from the environment. **Memory** — the ability to store and recall information across tasks. **Tools** — the systems the agent can use to take action, such as email, web search, or databases. And **Planning** — the ability to break a large goal into smaller, manageable steps.

All five components work together. Remove any one of them and the agent becomes significantly less effective. Together, they create a system capable of handling genuinely complex work.

SLIDE 8 — Applications: Where AI Agents Are Used Today

Person 1:
Where are AI Agents being used right now? Let me share some real examples — these are not future plans, they are live deployments.

Customer Support. Banks and retailers use AI agents that operate around the clock. They understand customer queries, check order status, process refunds, and escalate to a human only when necessary. A customer contacting support at midnight receives an immediate, accurate response — no waiting.

Software Development. Tools like GitHub Copilot do not just suggest a line of code — they understand the entire project, write complete functions, identify bugs, and run automated tests.

Financial Trading. In capital markets, AI agents analyse millions of data points per second and execute trades in milliseconds — a speed and precision no human trader can match.

Research. Instead of an analyst spending three days reading through papers, an AI research agent processes hundreds of documents overnight and delivers a structured summary the next morning.

Travel Planning. Agents that simultaneously connect to airlines, hotels, and car hire services to build a complete itinerary within seconds.

Cybersecurity. Security agents monitor networks continuously. The moment they detect an anomaly, they take action — blocking a threat before any human analyst has even seen the alert.

In every case, these agents are doing substantive work in the real world, not in a controlled test environment.

> ***AUDIENCE QUESTION Looking at these applications — which one surprises you most? And can anyone think of a repetitive or time-consuming task in your own work that an AI Agent could potentially handle? [Pause for two or three responses. This is a strong engagement moment — most people have an immediate answer.]***

SLIDE 9 — Data: AI Agent Adoption by Industry (Gartner 2024–2025)

Person 1:
This chart from Gartner's 2024–2025 report illustrates something important. The industries adopting AI Agents most rapidly are not limited to technology companies. Healthcare, retail, financial services, and manufacturing are all investing heavily.

The key point here is that this is not a niche technology for specialist teams. It is a broad shift happening across the entire economy. If your organisation operates in any of these sectors, AI Agents are either already part of your competitive landscape or will be very soon.

SLIDE 10 — Tools: AI Agents Use Tools Like Humans

Person 1:
Here is an idea that makes AI Agents much easier to understand: they use tools, just as we do.

Think about how you work. You use a search engine to find information. You use email to communicate. You use a spreadsheet to organise data. You do not do everything from memory — you use the right tool for each task.

AI Agents work in exactly the same way. They have access to a set of tools: web search, code execution, database access, APIs that connect to external systems, email and messaging services, and file storage. The agent decides which tool to use and at what point in the task — based on what the goal requires.

A useful analogy: think of a surgeon in an operating theatre. They have a full set of instruments — scalpels, clamps, sutures — each designed for a specific purpose. The surgeon's training tells them which instrument to pick and when. The AI Agent's brain does the same — it selects the right tool for each step of the task.

SLIDE 11 — Examples: Popular AI Agents in Use Today

Person 1:
Before I hand over to my colleague, let us look at some well-known AI Agents that are already available.

AutoGPT was one of the first agents made available to the public. You give it a goal — for example, 'research the five best electric cars and write a comparison report' — and it carries out the entire task independently.

Devin is the world's first AI software engineer. Given a brief description of an application, it writes the code, sets up the infrastructure, and deploys the product — without a human developer involved in any individual step.

Microsoft Copilot, built into Office 365, drafts emails, creates presentations from notes, summarises long documents, and analyses data in spreadsheets — operating as an agent inside tools your teams already use.

GPT Actions connects ChatGPT to applications such as Gmail, Slack, and databases. You can instruct it to find the latest invoice from a specific supplier, update a spreadsheet, and send a confirmation message — all in one instruction.

These are production tools being used by real organisations today.

That concludes the first section. I will now hand over to [Name 2], who will cover multi-agent systems, the business case, and the key considerations for deployment.

PERSON 2 — SLIDES 12 to 18 | Multi-Agent Systems · Benefits · Challenges · How to Build

SLIDE 12 — Advanced: Multi-Agent Systems: Teamwork

Person 2:
Thank you, [Name 1]. Hello, everyone. I am [Your Name], and I will be covering the next section.

We have established what a single AI Agent can do. Now let us look at what happens when multiple agents work together — because that is where the real power lies.

Consider how a hospital emergency department operates. When a patient arrives, one person does not handle everything. A triage nurse assesses priority, a doctor makes the clinical decision, a radiologist reads the scans, a pharmacist prepares medication, and an administrator manages the records — all simultaneously, all communicating with each other. Together, they achieve an outcome that no single individual could deliver as quickly or accurately alone.

A Multi-Agent System works on the same principle. A Manager Agent receives the overall task and distributes the work. Specialist agents handle specific parts: a Research Agent gathers information, a Writing Agent produces content, a Coder Agent builds what is needed, and an Email Agent manages communication. They all operate in parallel and feed their outputs back to the Manager Agent.

The result: what might take a human team an entire working day can be completed in minutes — because multiple agents are working simultaneously, each expert in its own area.

AUDIENCE QUESTION A practical question for the room: if you could deploy a team of AI agents within your own team at work, what roles would you assign them? What would the Manager Agent delegate first? [Allow two or three responses. This makes the concept directly relevant to the audience.]

SLIDE 13 — Benefits: Why AI Agents Are Powerful

Person 2:
Let us now look at why organisations are investing in this technology. What does it actually deliver in practice?

Speed. AI Agents complete tasks approximately one hundred times faster than a person working through the same process. A task that takes a team member three hours — gathering data, organising it, producing a report — takes an agent under two minutes. Not because it cuts corners, but because it does not pause, does not switch between tasks, and does not need any breaks.

Accuracy. Repetitive tasks are where human error is most likely. Fatigue, distraction, and volume all affect quality over time. AI Agents do not experience any of those factors. A quality-checking agent running at eleven at night performs identically to one running at nine in the morning.

Scale. One AI Agent can handle a thousand tasks simultaneously. That is not an estimate — it is a structural capability. One system, many parallel workstreams, at the same time.

Cost reduction. Over time, automating repetitive, high-volume tasks reduces operational costs by approximately sixty percent. The initial investment pays for itself relatively quickly.

Continuous availability. AI Agents do not have weekends, annual leave, or sick days. A security monitoring agent operates with the same level of attention at three on a Sunday morning as it does on a Tuesday afternoon.

Adaptability. Unlike traditional rule-based software, which breaks when conditions change, AI Agents learn and adjust. They improve with experience.

In combination, these capabilities allow organisations to deliver outcomes that were previously not achievable — at a speed and quality that traditional processes simply cannot match.

SLIDE 14 — Challenges: Challenges and Limitations

Person 2:

It is equally important to be clear-eyed about the limitations. AI Agents are powerful, but they are not without risk. Understanding the challenges is what allows you to deploy the technology responsibly.

Hallucinations. This is the term used when an AI produces incorrect information — and presents it with confidence. An agent might include a statistic that does not exist or cite a source that is inaccurate. For high-stakes decisions, human review of the agent's output remains essential.

Security risks. An agent with access to email, files, and databases is also a potential attack surface. A technique called prompt injection allows a bad actor to embed hidden instructions within content the agent reads — essentially tricking it into performing unintended actions. Think of it as a forged note inside a delivery that redirects the courier to a different address. The courier follows the note, unaware it was not legitimate.

Operating cost. Every decision an AI Agent makes involves calling a large language model, which carries a computational cost. Poorly designed agents that make unnecessary calls can become expensive at scale. Efficient design matters.

Over-automation risk. If an agent makes an error early in a task and there are no checkpoints, that error can propagate through many subsequent steps before anyone notices. Human oversight at key decision points is not optional — it is a design requirement.

Limited common sense. AI Agents can solve highly complex technical problems, yet may fail at tasks that require basic real-world judgement. They do not have lived experience or intuition.

Privacy and compliance. An agent that accesses personal data must be governed carefully. Questions of data retention, access controls, and regulatory compliance must be addressed before any agent is deployed in a professional environment.

The message is not to avoid the technology. It is to use it with the right guardrails in place. With proper design and oversight, all of these challenges are manageable.

AUDIENCE QUESTION Of these six challenges, which concerns you most — in the context of your own work or organisation? And do you have thoughts on how it could be addressed? [Allow three or four responses. Security and hallucinations typically generate the most discussion.]

SLIDE 15 — Market Data: The AI Agent Market

Person 2:
Let us put some scale on this. The AI Agent market is projected to reach 102 billion dollars by 2028, growing at 45 percent per year.

For context: the global smartphone market — built over fifteen years — is currently valued at around 500 billion dollars. AI Agents are on track to reach a quarter of that size in just four years. That is an extraordinary pace of adoption.

Two factors are driving this. First, cost efficiency — organisations are under consistent pressure to do more with less, and AI Agents address that directly. Second, the technology has matured. A few years ago, early agents were inconsistent and required significant human supervision. Today, they are reliable enough for real business workloads.

If you are responsible for strategy or planning in your organisation, this number carries a clear implication: organisations that develop AI Agent capabilities now will hold a structural advantage over those that delay. This is not a technology to monitor from a distance.

SLIDE 16 — Build: How to Build a Simple AI Agent

Person 2:
For those interested in how an AI Agent is actually built, this slide walks through the six key steps.

Step 1 — Define the goal. Be precise. 'Handle customer emails' is not specific enough. 'Respond to customer delivery queries within two hours, using the shipping database to provide accurate status updates' — that is a well-defined goal.

Step 2 — Choose the AI model. Select the language model that will power the agent's reasoning — GPT-4, Claude, Gemini, or another. Each has different strengths depending on the task.

Step 3 — Connect the tools. Determine what the agent needs to interact with — a database, an email service, a web search API — and connect it to those systems.

Step 4 — Add memory. Decide whether the agent needs to retain context across interactions. A customer service agent that forgets the start of a conversation mid-way through is not effective.

Step 5 — Set boundaries. Define what the agent is not permitted to do, and which actions require human approval before execution. Every agent needs clear operational limits.

Step 6 — Test and refine. Run the agent through realistic scenarios. Identify where it fails. Make adjustments. Repeat.

You do not need to be a software developer to understand this process. It follows the same logic as onboarding a new team member — you define the role, provide the right tools and resources, set the rules, and then review their work as they get up to speed.

SLIDE 17 — Frameworks: Popular AI Agent Frameworks

Person 2:
For the more technically-minded in the room, this slide covers the main frameworks — the development toolkits — that engineers use to build AI Agents.

LangChain is the most widely used starting point. It simplifies the process of connecting a language model to memory and tools, and is well-suited to initial projects.

LangGraph extends LangChain for more complex workflows, where an agent may need to loop back, branch, or handle conditional logic.

AutoGen, from Microsoft, is designed for multi-agent setups — systems where multiple agents communicate with each other to collaboratively solve a problem.

CrewAI allows you to define a team of agents with distinct roles, goals, and areas of responsibility — similar to designing an organisational structure.

Semantic Kernel, also from Microsoft, is an enterprise-grade framework designed for large-scale production deployments.

If you want to explore this practically, LangChain has clear documentation and tutorials online. It is possible to have a basic working agent running within a single afternoon.

I will now hand over to [Name 3], who will take us through the future of AI Agents, two major real-world case studies, and close out our session.

PERSON 3 — SLIDES 18 to 24 | Future · LLMs · RAG · Case Studies · Closing

SLIDE 18 — Future: The Future of AI Agents (Timeline)

Person 3:
Thank you, [Name 2]. Hello, everyone. I am [Your Name], and I will be covering the final part of our session.

Let us start with a timeline, because it helps to understand not just where this technology is going, but how quickly.

In **2024** — already behind us — agents became capable of browsing the web, writing and running code, and reading files. That was the year agents moved from experimental to genuinely useful.

In **2025** — right now — we are in the era of agentic workflows. Agents handle long, multi-step tasks that run for hours or days with minimal human involvement. A team might initiate an agent on a Monday and receive a completed research report by Wednesday.

By **2026 to 2027**, physical AI Agents — robots controlled by AI agent systems — will become mainstream in warehouses, hospitals, and construction sites. Early versions are already operational with companies like Boston Dynamics and in Amazon's fulfilment centres.

By **2028 to 2030**, the projections point to society-scale applications — agents managing traffic systems, energy grids, and healthcare logistics. These are direct extensions of what is being built today.

Sam Altman, CEO of OpenAI, has stated that AI Agents will do the work of a small company. The implication is significant: the role of people will shift from doing tasks directly to managing the agents that perform them. That is a meaningful change in how organisations operate.

AUDIENCE QUESTION Looking at this timeline honestly — what is your reaction? Excited, cautious, or a combination of both? [Pause. Acknowledge all responses without judgement — most people feel a mix of both, and that is a completely reasonable position. Both reactions are valid. The key is to be informed and prepared, which is exactly what today is about.]

SLIDE 20 — LLM Brain: The Role of Large Language Models in AI Agents

Person 3:
Let me spend a moment on the technology that sits at the heart of every AI Agent: the Large Language Model, or LLM.

An LLM is what separates an AI Agent from basic automation. Traditional automation operates on fixed rules — if this happens, do that. An LLM can understand context, interpret nuance, and reason through a problem.

A concrete example: a traditional automated system for customer emails might be programmed to respond to the word 'refund.' But if a customer writes 'I ordered this three weeks ago and I am still waiting,' a rule-based system may not recognise that as a refund situation at all. An LLM-powered agent understands the frustration, identifies the likely need, and responds appropriately — without being given a specific keyword to match.

The LLM also coordinates the agent's tools. It determines: for this step, I need to search the web. For this step, I need to retrieve a database record. For this step, I need to send a notification. It is the central decision-maker that orchestrates the entire process.

Without an LLM, you have automation. With an LLM, you have intelligence. That distinction matters enormously in practice.

SLIDE 21 — RAG: Retrieval-Augmented Generation

Person 3:
Let me introduce a concept that sounds technical but is straightforward once you see how it works: RAG, or Retrieval-Augmented Generation.

The core challenge with AI models is that they were trained on data up to a fixed date. They do not automatically know your organisation's internal policies, your current product catalogue, or this morning's news.

RAG solves this problem. Before the agent responds, it first searches a knowledge base — your company documents, a database, a library of records — retrieves the most relevant information, and uses that as the basis for its response.

Here is a useful way to think about it. Imagine you have just joined an organisation and a colleague asks you a specific question about a product. You do not know the answer from memory, but you know exactly where to find it. You go to the product documentation, read the relevant section, and answer accurately and confidently. That is RAG. The agent retrieves the right source before responding — rather than guessing.

This is what makes AI Agents reliable in professional and regulated environments. A medical AI that queries actual patient records. A legal AI that reviews actual case files. A customer service agent that references current product documentation. Each is grounded in verified information, not inference.

AUDIENCE QUESTION Can anyone think of a situation in your own work where an AI that could instantly search your organisation's internal knowledge base and give you a precise, accurate answer would save meaningful time? [Pause for two or three responses. Common examples include onboarding, compliance queries, and policy lookups.]

SLIDE 22 — Case Study: Netflix — AI Agents at Scale

Person 3:
Let us look at a company most people in this room will be familiar with: Netflix.

Netflix has over 250 million subscribers across 190 countries. Every single one of them sees a different home screen when they log in. That level of personalisation at that scale is only possible through AI Agents.

Netflix's recommendation system is an agent that analyses what you watch, what you pause or rewind, what you search for, and what people with similar taste profiles have watched. It uses all of that to decide which titles to present to you, in what order, and with which thumbnail image — for every individual session.

Beyond recommendations, AI agents at Netflix automatically test video playback across more than fifty device types, predict which subscribers are most likely to cancel in the coming month and trigger personalised retention campaigns before they do, and continuously optimise server infrastructure to manage costs.

The measurable results: eighty percent of viewing on Netflix is driven by agent recommendations rather than direct browsing. Their subscriber retention rate is forty percent better than the industry average. And they save over one billion dollars per year in infrastructure efficiency.

This is not a pilot project. This is AI Agents operating at enormous scale, delivering concrete business outcomes every single day.

SLIDE 23 — Case Study: Tesla — AI Agents in Autonomous Vehicles

Person 3:
Our second case study is Tesla — perhaps the most technically complex real-world deployment of AI Agents in existence.

When a Tesla operates in self-driving mode, it is not running a single AI system. It is running a coordinated team of AI Agents in real time, each with a specific responsibility.

The **Vision Agent** processes footage from eight cameras simultaneously, analysing 2.5 billion frames per day. It identifies pedestrians, cyclists, road markings, and traffic signals in real time with 99.9 percent accuracy.

The **Decision Agent** takes that visual data and determines the vehicle's next action — accelerate, brake, change lane. It plans five to ten seconds ahead at every moment.

The **Safety Agent** monitors the vehicle's systems — battery, brakes, temperature — and alerts the driver to developing issues before they become critical.

The **Fleet Learning Agent** is particularly notable. There are 1.5 million Tesla vehicles on the road globally. Every unusual driving event any of them encounters is anonymously shared and used to improve the system for all vehicles. It is the equivalent of 1.5 million drivers sharing their collective road experience every single day. No conventional training programme could achieve that level of scale.

The outcomes: 99.99 percent incident reduction over time, full self-driving capability in active beta, and agents processing 2.5 billion visual frames daily. This is multi-agent collaboration operating at a level of real-time precision and scale that is simply not achievable any other way.

SLIDE 19 — Summary: Key Takeaways

Person 3:
Before we close, let me bring together the key points from today's session.

First: An AI Agent is a software program that can perceive its environment, reason through a problem, take action, and learn from the outcome — without step-by-step human instruction.

Second: Agents use tools — web search, databases, APIs, communication systems — and combine them intelligently to solve complex, multi-step problems.

Third: This technology is not confined to the technology sector. Every major industry is deploying it now.

Fourth: The measurable benefits are speed, accuracy, scalability, cost reduction, and continuous availability.

Fifth: The challenges — hallucinations, security, cost, and privacy — are real. They are manageable with the right design, but they must be taken seriously.

Sixth: The market will reach 102 billion dollars by 2028. The window for building early capability and competitive advantage is now, not later.

AUDIENCE QUESTION OPEN Q&A; — We have time for questions. Is there anything from today's session you would like to explore further, or anything you would like clarified? [Allow 10 to 15 minutes. All three presenters should feel comfortable contributing answers.]

SLIDE 24 — Thank You and Closing

Person 3:
We are close to time, so let me close with a few final thoughts.

What we have discussed today represents a genuine shift in how work gets done — not a distant possibility, but something that is already happening across industries and organisations right now.

You do not need to be a developer to benefit from this. You need to understand what is possible, start identifying where AI Agents could add value in your context, and begin building familiarity with the tools.

Three practical steps to take from here: continue learning — follow developments in this space and try available tools. Start experimenting — even a simple test with LangChain or AutoGPT will teach you more than reading about it. And stay current — this field moves quickly, and staying informed is itself a significant advantage.

On behalf of [Name 1], [Name 2], and myself — thank you very much for your time and your engagement today. We hope this session has been genuinely useful. We are happy to continue any of these conversations directly.

Thank you.

End of script | *Estimated total duration: 55 – 65 minutes including Q&A pauses* | *If audience discussion is rich, allow it to run — engagement is more valuable than strict timing.*